



From Simple Models to Topic Models

- get a big picture of textual data



Outline

- Create Simple Text Model for Classification
 - fitcecoc
- Analyze Text Data Using Multiword Phrases
 - bagOfNgrams
- Analyze Text Data Using Topic Models
 - fitlda



Create Simple Text Model for Classification

- We show how to train a simple text classifier on word frequency counts using a bag-of-words model.
- We create a simple classification model which uses word frequency counts as predictors. This example trains a simple classification model to predict the category of factory reports using text descriptions.



Load and Extract Text Data

- Load the example data. The file `factoryReports.csv` contains factory reports, including a text description and categorical labels for each report.

```
filename = "factoryReports.csv";  
data = readtable(filename,'TextType','string');  
head(data)
```



Load and Extract Text Data

head(data)

ans =

8×5 [table](#)

Description	Category	Urgency	Resolution	Cost
"Items are occasionally getting stuck in the scanner spools."	"Mechanical Failure"	"Medium"	"Readjust Machine"	45
"Loud rattling and banging sounds are coming from assembler pistons."	"Mechanical Failure"	"Medium"	"Readjust Machine"	35
"There are cuts to the power when starting the plant."	"Electronic Failure"	"High"	"Full Replacement"	16200
"Fried capacitors in the assembler."	"Electronic Failure"	"High"	"Replace Components"	352
"Mixer tripped the fuses."	"Electronic Failure"	"Low"	"Add to Watch List"	55
"Burst pipe in the constructing agent is spraying coolant."	"Leak"	"High"	"Replace Components"	371
"A fuse is blown in the mixer."	"Electronic Failure"	"Low"	"Replace Components"	441
"Things continue to tumble off of the belt."	"Mechanical Failure"	"Low"	"Readjust Machine"	38

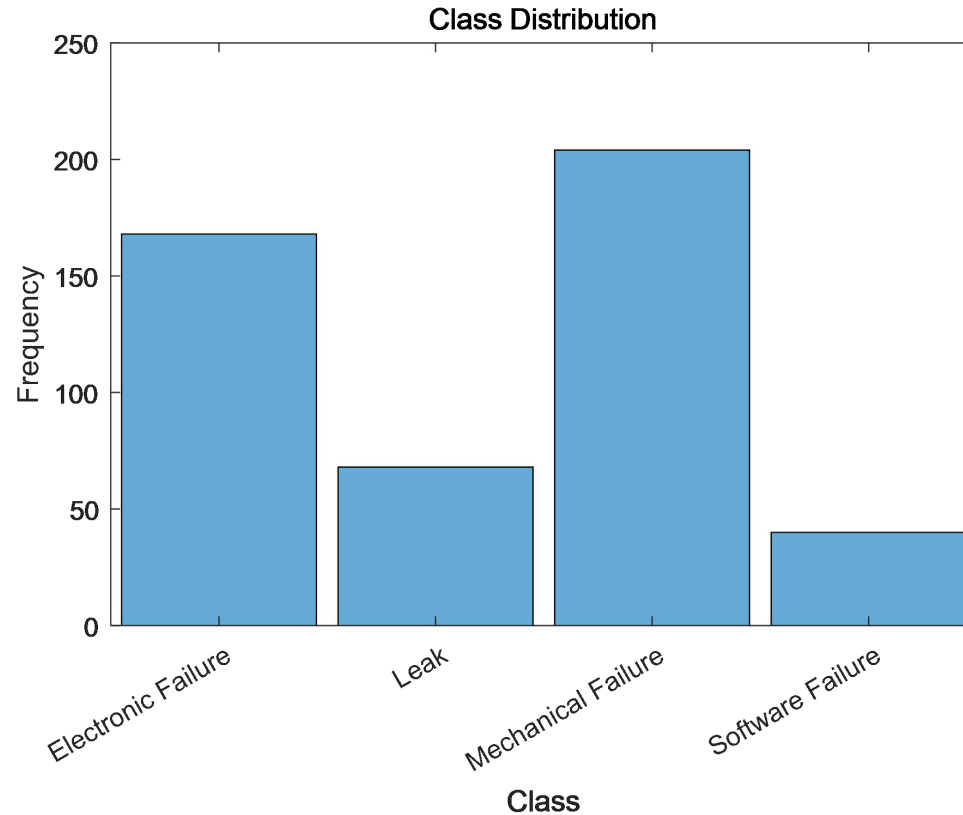


Load and Extract Text Data

- Convert the labels in the Category column of the table to categorical and view the distribution of the classes in the data using a histogram.

```
data.Category = categorical(data.Category);  
figure; histogram(data.Category)  
xlabel("Class")  
ylabel("Frequency")  
title("Class Distribution")
```

Load and Extract Text Data





Load and Extract Text Data

- Partition the data into a training partition and a held-out test set. Specify the holdout percentage to be 10%.

```
cvp = cvpartition(data.Category,'Holdout',0.1);  
dataTrain = data(cvp.training,:);  
dataTest = data(cvp.test,:);
```




Load and Extract Text Data

- Extract the text data and labels from the tables.

```
textDataTrain = dataTrain.Description;
```

```
textDataTest = dataTest.Description;
```

```
YTrain = dataTrain.Category;
```

```
YTest = dataTest.Category;
```



Prepare Text Data for Analysis

- Create a function which tokenizes and preprocesses the text data so it can be used for analysis. The function `preprocessText`, performs the following steps in order:
 - 1 Tokenize the text using `tokenizedDocument`.
 - 2 Remove a list of stop words ("and", "of") using `removeStopWords`.
 - 3 Lemmatize the words using `normalizeWords`.
 - 4 Erase punctuation using `erasePunctuation`.
 - 5 Remove words with 2 or fewer characters and with 15 or more characters using `removeShortWords` and `removeLongWords`.



Prepare Text Data for Analysis

- Use the example preprocessing function preprocessText to prepare the text data.

```
documents = preprocessText(textDataTrain);
```

```
documents(1:5)
```

```
ans =
```

5 × 1 tokenizedDocument:

6 tokens: items occasionally get stuck scanner spool

7 tokens: loud rattle bang sound come assembler piston

4 tokens: cut power start plant

3 tokens: fry capacitor assembler

3 tokens: mixer trip fuse



Prepare Text Data for Analysis

- Create a bag-of-words model from the tokenized documents.

```
bag = bagOfWords(documents)
```

```
bag =
```

```
bagOfWords with properties:
```

```
Counts: [432 × 336 double]
```

```
Vocabulary: [1 × 336 string]
```

```
NumWords: 336
```

```
NumDocuments: 432
```



Prepare Text Data for Analysis

- Remove words from the bag-of-words model that do not appear more than two times in total. Remove any documents containing no words from the bag-of-words model, and remove the corresponding entries in labels.

```
bag = removeInfrequentWords(bag,2);
```

```
[bag,idx] = removeEmptyDocuments(bag);
```

```
YTrain(idx) = [];
```



Train Supervised Classifier

- Train a supervised classification model using the word frequency counts from the bag-of-words model and the labels.
- Train a multiclass linear classification model using `fitcecoc`. Specify the `Counts` property of the bag-of-words model to be the predictors, and the event type labels to be the response. Specify the learners to be linear. These learners support sparse data input.

```
XTrain = bag.Counts;
```

```
mdl = fitcecoc(XTrain,YTrain,'Learners','linear')
```

Error-Correcting Output Codes

Class	Code Word														
	f_0	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_{10}	f_{11}	f_{12}	f_{13}	f_{14}
0	1	1	0	0	0	0	1	0	1	0	0	1	1	0	1
1	0	0	1	1	1	1	0	1	0	1	1	0	0	1	0
2	1	0	0	1	0	0	0	1	1	1	1	0	1	0	1
3	0	0	1	1	0	1	1	1	0	0	0	0	1	0	1
4	1	1	1	0	1	0	1	1	0	0	1	0	0	0	1
5	0	1	0	0	1	1	0	1	1	1	0	0	0	0	1
6	1	0	1	1	1	0	0	0	0	1	0	1	0	0	1
7	0	0	0	1	1	1	1	0	1	0	1	1	0	0	1
8	1	1	0	1	0	1	1	0	0	1	0	0	0	1	1
9	0	1	1	1	0	0	0	0	1	0	1	0	0	1	1

Figure 1: A 15 bit error-correcting output code for a ten-class problem



Test Classifier

- Predict the labels of the test data using the trained model and calculate the classification accuracy. The classification accuracy is the proportion of the labels that the model predicts correctly.
- Preprocess the test data using the same preprocessing steps as the training data. Encode the resulting test documents as a matrix of word frequency counts according to the bag-of-words model.

```
documentsTest = preprocessText(textDataTest);
```

```
XTest = encode(bag,documentsTest);
```




Test Classifier

- Predict the labels of the test data using the trained model and calculate the classification accuracy.

```
YPred = predict mdl, XTest;
```

```
acc = sum(YPred == YTest)/numel(YTest)
```

```
acc =
```

```
0.8958
```



Predict Using New Data

- Classify the event type of new factory reports.

str = ["Coolant is pooling underneath sorter."

"Sorter blows fuses at start up."

"There are very loud rattling sounds coming from the assembler."];

documentsNew = preprocessText(str);

XNew = encode(bag,documentsNew);

labelsNew = predict(mdl,XNew)

labelsNew = 3×1 categorical

Leak

Electronic Failure

Mechanical Failure



Analyze Text Data Using Multiword Phrases

- This example shows how to analyze text using n-gram frequency counts.
- An n-gram is a tuple of n consecutive words. For example, a bigram (the case when $n = 2$) is a pair of consecutive words such as "heavy rainfall". A unigram (the case when $n = 1$) is a single word. A bag-of-n-grams model records the number of times that different n-grams appear in document collections.



Analyze Text Data Using Multiword Phrases

- Using a bag-of-n-grams model, you can retain more information on word ordering in the original text data. For example, a bag-of-n-grams model is better suited for capturing short phrases which appear in the text, such as "heavy rainfall" and "thunderstorm winds".
- To create a bag-of-n-grams model, use **bagOfNgrams**. You can input bagOfNgrams objects into other Text Analytics Toolbox functions such as wordcloud and fitlda.



Load and Extract Text Data

- Load the example data. The file `factoryReports.csv` contains factory reports, including a text description and categorical labels for each event. Remove the rows with empty reports.

```
filename = "factoryReports.csv";  
data = readtable(filename,'TextType','String');
```



Load and Extract Text Data

- Extract the text data from the table and view the first few reports.

```
textData = data.Description;
```

```
textData(1:5)
```

```
ans = 5 × 1 string
```

```
"Items are occasionally getting stuck in the scanner spools."
```

```
"Loud rattling and banging sounds are coming from assembler pistons."
```

```
"There are cuts to the power when starting the plant."
```

```
"Fried capacitors in the assembler."
```

```
"Mixer tripped the fuses."
```



Load and Extract Text Data

- Use the example preprocessing function `preprocessText` to prepare the text data.

```
documents = preprocessText(textData);
```

```
documents(1:5)
```

```
ans =
```

5 × 1 tokenizedDocument:

6 tokens: item occasionally get stuck scanner spool

7 tokens: loud rattle bang sound come assembler piston

4 tokens: cut power start plant

3 tokens: fry capacitor assembler ...



Create Word Cloud of Bigrams

- Create a word cloud of bigrams by first creating a bag-of-n-grams model using `bagOfNgrams`.

```
bag = bagOfNgrams(documents)
```

```
bag =
```

```
bagOfNgrams with properties:
```

```
Counts: [480 × 941 double]
```

```
Vocabulary: [1 × 351 string]
```

```
Ngrams: [941 × 2 string]
```

```
NgramLengths: 2
```

```
NumNgrams: 941    NumDocuments: 480
```




Fit Topic Model to Bag-of-N-Grams

- A Latent Dirichlet Allocation (LDA) model is a topic model which discovers underlying topics in a collection of documents and infers the word probabilities in topics.
- Create an LDA topic model with 10 topics using `fitlda`. The function fits an LDA model by treating the n-grams as single words.

```
mdl = fitlda(bag,10,'Verbose',0);
```



More About Latent Dirichlet Allocation

- A latent Dirichlet allocation (LDA) model is a document topic model which discovers underlying topics in a collection of documents and infers word probabilities in topics. LDA models a collection of D documents as topic mixtures $\theta_1, \dots, \theta_D$, over K topics characterized by vectors of word probabilities ϕ_1, \dots, ϕ_K . It assumes that the topic mixtures $\theta_1, \dots, \theta_D$, and the topics ϕ_1, \dots, ϕ_K follow a Dirichlet distribution with concentration parameters α and β respectively.



More About Latent Dirichlet Allocation

- The topic mixtures $\theta_1, \dots, \theta_D$ are probability vectors of length K , where K is the number of topics. The entry θ_{di} is the probability of topic i appearing in the d th document.
- The topics ϕ_1, \dots, ϕ_K are probability vectors of length V , where V is the number of words in the vocabulary. The entry ϕ_{iv} corresponds to the probability of the v th word of the vocabulary appearing in the i th topic.



More About Latent Dirichlet Allocation

- Given the topics ϕ_1, \dots, ϕ_K and Dirichlet prior α on the topic mixtures, LDA assumes the following generative process for a document:

1 Sample a topic mixture $\theta \sim \text{Dirichlet}(\alpha)$. The random variable θ is a probability vector of length K , where K is the number of topics.



More About Latent Dirichlet Allocation

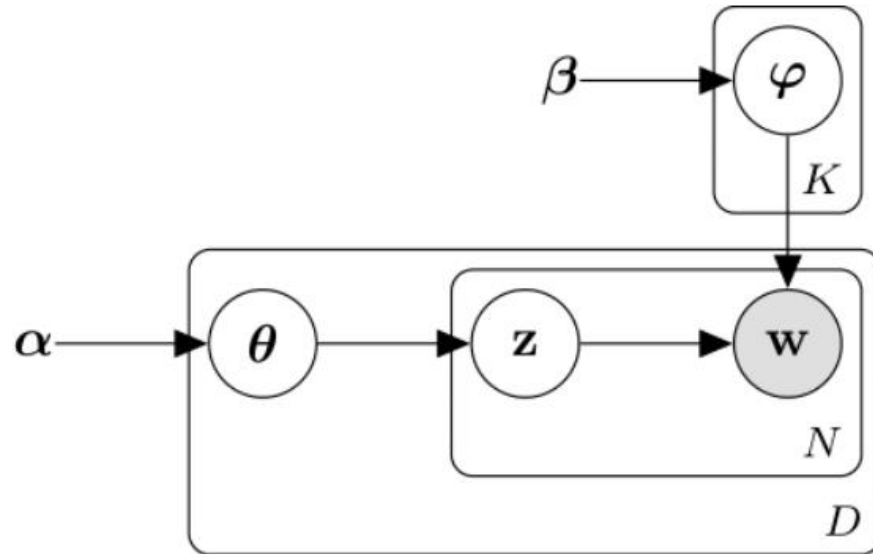
2 For each word in the document:

a) Sample a topic index $z \sim \text{Categorical}(\theta)$. The random variable z is an integer from 1 through K , where K is the number of topics.

b) Sample a word $w \sim \text{Categorical}(\phi_z)$. The random variable w is an integer from 1 through V , where V is the number of words in the vocabulary, and represents the corresponding word in the vocabulary.

More About Latent Dirichlet Allocation

LDA model is a probabilistic graphical model. The shaded nodes are observed variables, unshaded nodes are latent variables, nodes without outlines are model parameters.





Fit Topic Model to Bag-of-N-Grams

- Visualize the first four topics as word clouds. The word clouds highlight commonly co-occurring bigrams in the LDA topics. The function plots the bigrams with sizes according to their probabilities for the specified LDA topics.

figure

for i = 1:4

 subplot(2,2,i)

 wordcloud(mdl,i);

 title("LDA Topic " + i)

end

LDA Topic 1



LDA Topic 2



LDA Topic 3



LDA Topic 4





Analyze Text Using Longer Phrases

- To analyze text using longer phrases, specify the 'NGramLengths' option in `bagOfNgrams` to be a larger value.
- When working with longer phrases, it can be useful to keep stop words in the model. For example, to detect the phrase "is not happy", keep the stop words "is" and "not" in the model.
- Preprocess the text. Erase the punctuation using `erasePunctuation`, and tokenize using `tokenizedDocument`.

```
cleanTextData = erasePunctuation(textData);
```

```
documents = tokenizedDocument(cleanTextData);
```



Analyze Text Using Longer Phrases

- To count the n-grams of length 3 (trigrams), use `bagOfNgrams` and specify '`NGramLengths`' to be 3.

```
bag = bagOfNgrams(documents,'NGramLengths',3);
```

- Visualize the bag-of-n-grams model using a word cloud. The word cloud of trigrams better shows the context of the individual words.

figure

```
wordcloud(bag);
```

```
title("Text Data: Trigrams")
```




Analyze Text Using Longer Phrases

- View the top 10 trigrams and their frequency counts using `topkngrams`.

```
tbl = topkngrams(bag,10)
```

```
tbl=10×3 table
```

Ngram Count	NgramLength
-------------	-------------

"in" "the" "mixer"	14 3
--------------------	------

"in" "the" "scanner"	13 3
----------------------	------

"blown" "in" "the"	9 3
--------------------	-----

"the" "robot" "arm"	7 3
---------------------	-----

"stuck" "in" "the"	6 3
--------------------	-----

"is" "spraying" "coolant"	6 3
---------------------------	-----

"from" "time" "to"	6 3
--------------------	-----

"time" "to" "time"	6 3
--------------------	-----

"heard" "in" "the"	6 3
--------------------	-----

"on" "the" "floor"	6 3
--------------------	-----



Analyze Text Data Using Topic Models

- This example shows how to use the Latent Dirichlet Allocation (LDA) topic model to analyze text data.
- A Latent Dirichlet Allocation (LDA) model is a topic model which discovers underlying topics in a collection of documents and infers the word probabilities in topics.



Load and Extract Text Data

- Load the example data. The file `factoryReports.csv` contains factory reports, including a text description and categorical labels for each event.

```
data = readtable("factoryReports.csv", 'TextType', 'string');
```

```
head(data)
```

```
ans=8×5 table
```

```
Description Category Urgency Resolution Cost
```

```
"Items are occasionally getting stuck in the scanner spools." "Mechanical Failure"  
"Medium" "Readjust Machine" 45
```

```
"Loud rattling and banging sounds are coming from assembler pistons." "Mechanical  
Failure" "Medium" "Readjust Machine" 35 ...
```




Load and Extract Text Data

- Extract the text data from the field Description.

```
textData = data.Description;
```

```
textData(1:5)
```

```
ans = 5 × 1 string
```

```
"Items are occasionally getting stuck in the scanner spools."
```

```
"Loud rattling and banging sounds are coming from assembler pistons."
```

```
"There are cuts to the power when starting the plant."
```

```
"Fried capacitors in the assembler."
```

```
"Mixer tripped the fuses."
```




Prepare Text Data for Analysis

- Use the preprocessing function `preprocessText` to prepare the text data.

```
documents = preprocessText(textData);
```

```
documents(1:5)
```

```
ans =
```

5 × 1 tokenizedDocument:

6 tokens: items occasionally get stuck scanner spool

7 tokens: loud rattle bang sound come assembler piston

4 tokens: cut power start plant

3 tokens: fry capacitor assembler

3 tokens: mixer trip fuse



Prepare Text Data for Analysis

- Create a bag-of-words model from the tokenized documents.

```
bag = bagOfWords(documents)
```

```
bag =
```

```
bagOfWords with properties:
```

```
Counts: [480 × 351 double]
```

```
Vocabulary: [1 × 351 string]
```

```
NumWords: 351
```

```
NumDocuments: 480
```



Prepare Text Data for Analysis

- Remove words from the bag-of-words model that have do not appear more than two times in total. Remove any documents containing no words from the bag-of-words model.

```
bag = removeInfrequentWords(bag,2);
```

```
bag = removeEmptyDocuments(bag)
```

```
bag =
```

```
bagOfWords with properties:
```

```
Counts: [480 × 162 double]
```

```
Vocabulary: [1 × 162 string]
```

```
NumWords: 162
```

```
NumDocuments: 480
```



Fit LDA Model

- Fit an LDA model with 7 topics.

```
numTopics = 7;
```

```
mdl = fitlda(bag,numTopics,'Verbose',0);
```

- If you have a large dataset, then the stochastic approximate variational Bayes solver is usually better suited as it can fit a good model in fewer passes of the data. The default solver for fitlda (collapsed Gibbs sampling) can be more accurate at the cost of taking longer to run. To use stochastic approximate variational Bayes, set the 'Solver' option to 'savb'.



Visualize Topics Using Word Clouds

- You can use word clouds to view the words with the highest probabilities in each topic. Visualize the first four topics using word clouds.

figure;

for topicIdx = 1:4

 subplot(2,2,topicIdx)

 wordcloud(mdl,topicIdx);

 title("Topic " + topicIdx)

end

[illegible]



View Mixtures of Topics in Documents

- Use transform to transform the documents into vectors of topic probabilities.

```
newDocument = tokenizedDocument("Coolant is pooling underneath sorter.");
```

```
topicMixture = transform mdl, newDocument);
```

```
figure
```

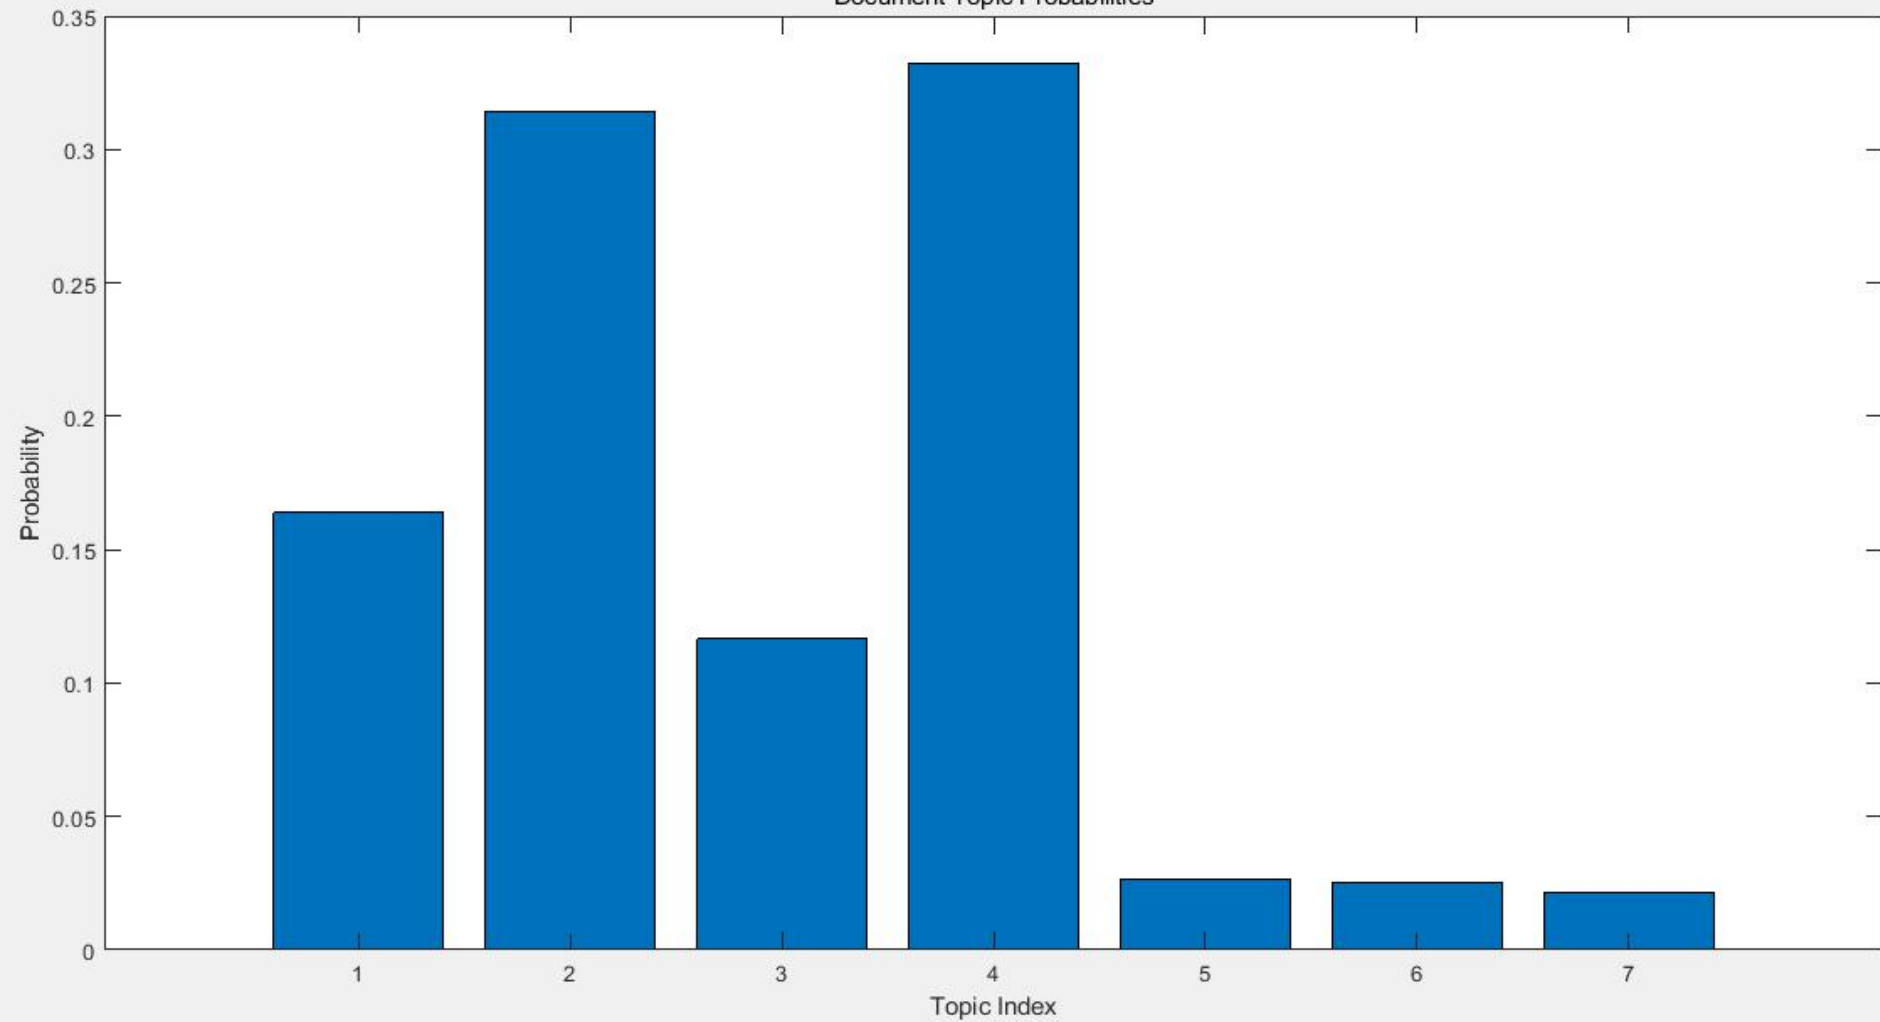
```
bar(topicMixture)
```

```
xlabel("Topic Index")
```

```
ylabel("Probability")
```

```
title("Document Topic Probabilities")
```

Document Topic Probabilities





View Mixtures of Topics in Documents

- Visualize multiple topic mixtures using stacked bar charts.
Visualize the topic mixtures of the first 5 input documents.

figure

```
topicMixtures = transform(mdl,documents(1:5));
```

```
barh(topicMixtures(1:5,:), 'stacked')
```

```
xlim([0 1])
```

```
title("Topic Mixtures")
```

```
xlabel("Topic Probability")
```

```
ylabel("Document")
```

```
legend("Topic " + string(1:numTopics), 'Location', 'northeastoutside')
```

